

ALTAF ANSARI

AI Engineer & Data Scientist

CONTACT

altafansari@pda@gmail.com
+91 9108778332
linkedin.com/in/altaf-ansari309
github.com/altaf-ansari

TECHNICAL SKILLS

Generative AI & LLMs

Agents • RAG • LangChain • LangGraph • OpenAI • Azure

ML / DL / NLP

PyTorch • TensorFlow • Scikit-learn • Transformers •

Backend & Deployment

Python • FastAPI • Docker • Kubernetes • AWS (EC2, ECR,

Data & Analytics

Pandas • NumPy • Spark • Hadoop • SQL Server • Power

Tools & Automation

Claude Code • GitHub Copilot • Cursor • n8n • Git •

CERTIFICATIONS

✓ AWS Certified ML — Specialty

Amazon Web Services

✓ Azure AI Fundamentals (AI-900)

Microsoft Azure

✓ Generative AI with LLMs

DeepLearning.AI

EDUCATION

B.E. Mechanical Engineering

Visvesvaraya Technological
University (VTU) — 2020

FOCUS AREAS

- Agentic AI Systems
- Multi-Agent Orchestration
- LLM Fine-tuning & Evaluation
- Cloud-Native AI Pipelines
- AI-Assisted Engineering
- Intelligent Document Processing

PROFESSIONAL SUMMARY

AI Engineer and Software Engineer with 3+ years of experience building production-grade AI systems, agentic workflows, and cloud-native ML pipelines. Specialized in LLM integration (RAG, agents, fine-tuning), FastAPI backend services, and end-to-end deployment on AWS and Azure. Track record of shipping intelligent products that improve accuracy by 30%+, reduce latency by 40%, and automate operations across enterprise environments.

WORK EXPERIENCE

Data Scientist

PharmaForce Infosoft Insights Pvt. Ltd | Dec 2024 — Present

Leading AI/ML initiatives across the 340B program — 9+ production systems shipped

- Built and deployed AI-powered backend services using FastAPI with LLM-based RAG pipelines, improving retrieval accuracy by 30% and reducing latency by 40%
- Designed agentic workflows using LangChain and LangGraph for automated data analysis, query generation, and multi-step reasoning across enterprise data
- Developed text-to-SQL APIs and semantic search services serving non-technical stakeholders, cutting data request turnaround from days to seconds
- Built cloud-native AI invoice extraction pipeline on Azure (Functions, Blob Storage, Azure OpenAI, PyMuPDF) processing 200+ invoices/month across multiple formats
- Fine-tuned LLMs locally and on AWS SageMaker for domain-specific extraction and classification tasks, improving baseline accuracy by 15%+
- Guided 3+ development teams on efficient adoption of Claude Code, reducing average feature delivery time
- Reverse-engineered legacy application versions and rebuilt them on modern Python/FastAPI architectures
- Built automated defect report newsletter pipeline using n8n, delivering weekly reports to 50+ stakeholders
- Developed multi-level agent orchestration POC on LangGraph with hierarchical sub-agent delegation for complex reasoning

Data Scientist

People Tech Group | May 2021 — May 2023

Built CV and ML systems handling 5K+ daily requests, reducing claims by 25%

- Developed and deployed YOLOv5/v8-based computer vision service achieving 94% accuracy, reducing manual review time by 70%
- Built ML-based risk prediction models for insurance, reducing claim costs by 25% across the portfolio
- Built and maintained end-to-end data pipelines for training and inference on AWS (EC2, ECR, S3)
- Deployed containerized ML services using Docker and Kubernetes ensuring high availability at scale
- Mentored 6 interns and conducted technical code reviews across multiple engineering teams

KEY PROJECTS

• Kareem Agent

LLM-based chatbot and agentic system automating support and data analysis workflows. Combines RAG, semantic search, and multi-step reasoning to handle complex enterprise queries autonomously.

• AI Invoice Extraction

Cloud-native Azure pipeline — auto-triggers on PDF upload, extracts structured data from diverse invoice formats via Azure OpenAI + PyMuPDF, and persists to SQL Server. Processes 200+ invoices/month.

• Text-to-SQL Platform

Natural language to SQL APIs enabling non-technical stakeholders to query business data directly. Built-in guardrails prevent destructive queries. Cut data request turnaround from days to seconds.

• Computer Vision Platform

YOLOv5/v8 automated detection system achieving 94% accuracy on production data, handling 5K+ daily requests and reducing manual review by 70%.

• Semantic Search & RAG

Hybrid retrieval combining dense vector embeddings with LLM synthesis for citation-backed enterprise search. Improved retrieval accuracy by 30% over keyword baselines.

• LangGraph Multi-Agent System

Hierarchical multi-level agent orchestration with specialized sub-agents for complex, multi-step reasoning. Modular design enables adding new capabilities without touching existing agents.